

# SIMD in C++: auto-vectorization in a nutshell

Anthony Boulmier

Université de Genève

*anthony.boulmier@unige.ch*

October 20, 2020

# Outline

- 1 Introduction
- 2 How to ?
- 3 Breaking the code!
- 4 Benchmarks
- 5 Conclusion

# Motivation


- CPU clock rate does not increase anymore

# Motivation

- CPU clock rate does not increase anymore
- More cores → more transistors !

# Motivation

- CPU clock rate does not increase anymore
- More cores  $\rightarrow$  more transistors !
- Or ? Do multiple instructions at the same time!



# What is SIMD ?

SIMD stands for: **S**ingle **I**nstruction **M**ultiple **D**ata

- Same operation on multiple values (think about element-wise vector operations)

Example: You want to sum the forces (XYZ) for 100K particles. Do 4, 8, or 16 summations at the same time instead of 1.

# What is SIMD ?

SIMD stands for: **S**ingle **I**nstruction **M**ultiple **D**ata

- Same operation on multiple values (think about element-wise vector operations)  
Example: You want to sum the forces (XYZ) for 100K particles. Do 4, 8, or 16 summations at the same time instead of 1.
- Recent CPUs have an instruction set that works on **vector registers** (`%xmm`, `%ymm`, and `%zmm` registers, note: `xmm` are also used for scalar operations)

# What is SIMD ?

SIMD stands for: **S**ingle **I**nstruction **M**ultiple **D**ata

- Same operation on multiple values (think about element-wise vector operations)  
Example: You want to sum the forces (XYZ) for 100K particles. Do 4, 8, or 16 summations at the same time instead of 1.
- Recent CPUs have an instruction set that works on **vector registers** (`%xmm`, `%ymm`, and `%zmm` registers, note: `xmm` are also used for scalar operations)
- Up to 16x performance improvement and it is almost free lunch !  
Codes that benefit from vector instructions are called vectorized code.



# Vector registers






Figure: Vector registers, source: Cornell

# What is SIMD ?

## SIMD Mode



## Scalar Mode



# A LOT can be done with SIMD




Figure: What is this?

# How to write auto-vectorized code?

Some portion of your codes may already benefits from it!

Compilers “auto-vectorize” under the following constraints:

- Can only vectorize for-loop with a fixed (known at the beginning) number of iterations

# How to write auto-vectorized code?

Some portion of your codes may already benefits from it!

Compilers “auto-vectorize” under the following constraints:

- Can only vectorize for-loop with a fixed (known at the beginning) number of iterations
- No dependency between array indices (iterations are independent)

# How to write auto-vectorized code?

Some portion of your codes may already benefits from it!

Compilers “auto-vectorize” under the following constraints:

- Can only vectorize for-loop with a fixed (known at the beginning) number of iterations
- No dependency between array indices (iterations are independent)
- No conditional branching

# How to write auto-vectorized code?

Some portion of your codes may already benefits from it!

Compilers “auto-vectorize” under the following constraints:

- Can only vectorize for-loop with a fixed (known at the beginning) number of iterations
- No dependency between array indices (iterations are independent)
- No conditional branching
- The compiler is sure that the performance will increase (use of cost function)

# How to write auto-vectorized code?

Some portion of your codes may already benefits from it!

Compilers “auto-vectorize” under the following constraints:

- Can only vectorize for-loop with a fixed (known at the beginning) number of iterations
- No dependency between array indices (iterations are independent)
- No conditional branching
- The compiler is sure that the performance will increase (use of cost function)
- Only vectorize inner loops

Rule of thumb: if you can't tell how to vectorize a code, neither can the compiler.



# Write and debug vectorized code

- 1 Write auto-vectorization friendly codes  
*Array of Structure (AoS) vs. Structure of Array (SoA)*
- 2 Use the proper compilation line  
Vectorization is enabled by default at *-O3*  
*gcc -O2 -ftree-vectorize -fopt-info-vec-{all,missed,optimized}*  
*clang -O2 -ftree-vectorize -Rpass-analysis=loop-vectorize*  
*-Rpass=loop-vectorize*
- 3 Look at assembly code  
*-S -o main.s*

## Example

Listing 1: AoS

```
const auto N = 100000;
struct Body {
    std::array<float, 2> v, p /*, ... */;
};
std::array<Body, N> bodies;
```

## Example

### Listing 2: SoA

```
const auto N = 100000;
struct Bodies {
    std::array<float, N> vx, px, vy, py /*, ... */;
};
Bodies bodies;
```

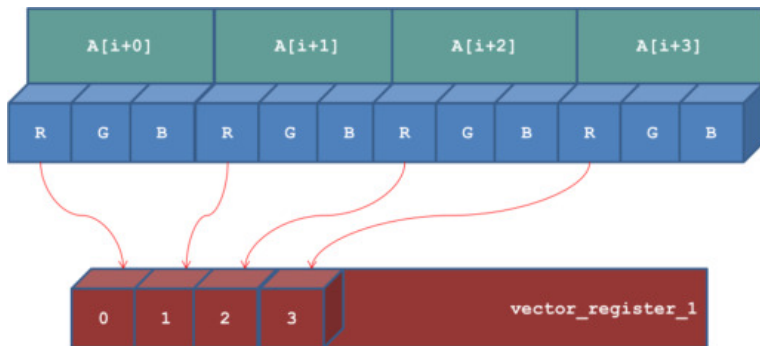


Figure: non-coalesced memory access, source: Jeffers et al. 2016

## Example

### Listing 3: Vectorization example

```
/* ... */  
float f(float * a, float * b){  
    #pragma GCC ivdep  
    for(int i=0; i < 10000000; ++i)  
        a[i] = a[i] + b[i];  
    return a[50000];  
}  
/* ... */
```

godbolt.org: online disassembler

## Behind the scene (non-vectorized)

### Listing 4: Vectorization example

```
.L2:  
    movss    xmm0, DWORD PTR [rdi+rax]  
    addss    xmm0, DWORD PTR [rsi+rax]  
    movss    DWORD PTR [rdi+rax], xmm0  
    add      rax, 4  
    cmp      rax, 400000000  
    jne      .L2
```

In [godbolt.org](https://godbolt.org): `gcc-10 -O2`

## Behind the scene (vectorized)

### Listing 5: Vectorization example

.L2:

```
vmovups zmm1, ZMMWORD PTR [rsi+rax]
vaddps  zmm0, zmm1, ZMMWORD PTR [rdi+rax]
vmovups ZMMWORD PTR [rdi+rax], zmm0
add     rax, 64
cmp     rax, 40000000
jne     .L2
```

In godbolt.org: `gcc-10 -O2 -ftree-vectorize -mavx512f`

## Behind the scene (vectorized)

Listing 6: Vectorization example

.L2:

```
vmovups zmm1, ZMMWORD PTR [rsi+rax]
vaddps  zmm0, zmm1, ZMMWORD PTR [rdi+rax]
vmovups ZMMWORD PTR [rdi+rax], zmm0
add     rax, 64
cmp     rax, 40000000
jne     .L2
```

In godbolt.org: `gcc-10 -O2 -ftree-vectorize -mavx512f`

What if the array size in byte is not divisible by the vector register size?



## Behind the scene (vectorized)

### Listing 7: Vectorization example

.L2:

```
vmovups zmm1, ZMMWORD PTR [rsi+rax]
vaddps  zmm0, zmm1, ZMMWORD PTR [rdi+rax]
vmovups ZMMWORD PTR [rdi+rax], zmm0
add     rax, 64
cmp     rax, 40000000
jne     .L2
```

In godbolt.org: `gcc-10 -O2 -ftree-vectorize -mavx512f`

What if the array size in byte is not divisible by the vector register size?

The compiler executes the rest in scalar mode.

Note: unaligned calls are equivalent to aligned code since *Intel Sandy Bridge*: see,

Agner Fog

## What does the compiler say?

Listing 8: Vectorization example

```
void f(float* A, float* B, int N){  
    for(int i = 0; i < N; ++i)  
        A[i] = A[i] * B[i];  
}
```

```
g++-10 main.cpp -O3 -mavx512f -fopt-info-vec-all -o main
```

## What does the compiler say?

### Listing 9: Vectorization example

```
void f(float* A, float* B, int N){  
    for(int i = 0; i < N; ++i)  
        A[i] = A[i] * B[i];  
}
```

```
g++-10 main.cpp -O3 -mavx512f -fopt-info-vec-all -o main
```

```
main.cpp:3:19: optimized: loop vectorized using 64 byte vectors  
main.cpp:3:19: optimized: loop versioned for vectorization  
because of possible aliasing
```

The compiler built a vectorized version, but it had to add aliasing checks

## What does the compiler say?

### Listing 10: Vectorization example

```
void f(float* A, float* B, int N){  
#pragma GCC ivdep  
    for(int i = 0; i < N; ++i)  
        A[i] = A[i] * B[i];  
}
```

```
gcc-10 main.cpp -O3 -mavx512f -fopt-info-vec-all -o main
```

## What does the compiler say?

### Listing 11: Vectorization example

```
void f(float* A, float* B, int N){  
#pragma GCC ivdep  
    for(int i = 0; i < N; ++i)  
        A[i] = A[i] * B[i];  
}
```

```
gcc-10 main.cpp -O3 -mavx512f -fopt-info-vec-all -o main
```

```
main.cpp:3:19: optimized: loop vectorized using 64 byte vectors
```

Compiled to avx512 instruction set !

## What does the compiler say?

### Listing 12: Vectorization example

```
float f(float* A, float* B, int N){  
#pragma GCC ivdep  
    for(int i = 1; i < N; ++i)  
        A[i] = A[i-1] * B[i-1];  
    return A[0];  
}
```

```
gcc-10 main.cpp -O3 -mavx512f -fopt-info-vec-all -o main
```

## What does the compiler say?

### Listing 13: Vectorization example

```
float f(float* A, float* B, int N){  
#pragma GCC ivdep  
    for(int i = 1; i < N; ++i)  
        A[i] = A[i-1] * B[i-1];  
    return A[0];  
}
```

```
gcc-10 main.cpp -O3 -mavx512f -fopt-info-vec-all -o main
```

```
main.cpp:5:18: missed: could not vectorize loop  
main.cpp:6:21: missed: not vectorized: no vectype  
for stmt: _8=*_7
```

## What does the compiler say?

### Listing 14: Vectorization example

```
float f(float* A, float* B, int N){  
#pragma GCC ivdep  
    for(int i = 1; i < N; ++i)  
        if(A[i] < 3) A[i] = A[i] * B[i];  
    return A[0];  
}
```

```
gcc-10 main.cpp -O3 -mavx512f -fopt-info-vec-all -o main
```



## What does the compiler say?

### Listing 15: Vectorization example

```
float f(float* A, float* B, int N){  
#pragma GCC ivdep  
    for(int i = 1; i < N; ++i)  
        if(A[i] < 3) A[i] = A[i] * B[i];  
    return A[0];  
}
```

```
gcc-10 main.cpp -O3 -mavx512f -fopt-info-vec-all -o main
```

```
main.cpp:5:18: missed: could not vectorize loop  
main.cpp:5:18: missed: not vectorized: control flow in loop.
```

# Benchmark description

- Two arrays  $(x,y)$  containing  $S = 10^7$  float (32 bits).
- Iterate over the arrays and execute a function that costs  $X$  flops. The function is generated and inlined at compile time using *constexpr*.
- Do this  $T$  times and compute min, max, and average computing time in serial mode and (auto)vectorized mode.

# Benchmark description

- Two arrays  $(x, y)$  containing  $S = 10^7$  float (32 bits).
- Iterate over the arrays and execute a function that costs  $X$  flops. The function is generated and inlined at compile time using *constexpr*.
- Do this  $T$  times and compute min, max, and average computing time in serial mode and (auto)vectorized mode.

- Function MULSUM:

$$f_{\text{MUL}}(X, x_i, y_i) = y_i * f_{\text{MUL}}(X - 1, y_i, x_i)$$

$$f_{\text{MUL}}(0, x_i, y_i) = y_i$$

$$\text{Example: } f(2, x_i, y_i) = y_i * x_i * y_i$$

# Benchmark description

- Two arrays  $(x, y)$  containing  $S = 10^7$  float (32 bits).
- Iterate over the arrays and execute a function that costs  $X$  flops. The function is generated and inlined at compile time using *constexpr*.
- Do this  $T$  times and compute min, max, and average computing time in serial mode and (auto)vectorized mode.

- Function MULSUM:

$$f_{\text{MUL}}(X, x_i, y_i) = y_i * f_{\text{MUL}}(X - 1, y_i, x_i)$$

$$f_{\text{MUL}}(0, x_i, y_i) = y_i$$

$$\text{Example: } f(2, x_i, y_i) = y_i * x_i * y_i$$

- Function FMASUM:

$$f_{\text{FMA}}(X, x_i, y_i) = y_i + \sum_1^X x_i y_i$$

$$\text{Example: } f_{\text{FMA}}(2, x_i, y_i) = y_i + x_i y_i + x_i y_i$$

Note: I consider *fused multiply add (FMA)* to cost 1 flop.

# Benchmark




Figure: Time, gcc-10, Intel i7-10750H (Comet Lake): with avx2 (ymm 256 bits)

# Benchmark




Figure: Speedup, gcc-10, Intel i7-10750H (Comet Lake): with avx2 (ymm 256 bits)

# Benchmark




Figure: Speedup, gcc-10, Intel i7-10750H (Comet Lake): with avx2 (ymm 256 bits)

main.cpp:69:3: remark: the cost-model indicates that interleaving is not beneficial [-Rpass-analysis=loop-vectorize]



# Benchmark




Figure: Speedup, clang-10, Intel i7-10750H (Comet Lake): with avx2 (ymm 256 bits)

# Benchmark summary

- ❶ Bad news 1: when there is a few FLOP per iteration, the computation is bounded by memory transfer. Push/Pull to/from vector registers can kill your perfs!

# Benchmark summary

- ❶ Bad news 1: when there is a few FLOP per iteration, the computation is bounded by memory transfer. Push/Pull to/from vector registers can kill your perfs!
- ❷ Bad news 2: the compiler can mispredict performance and it may not vectorize, even straightforward code.

# Benchmark summary

- ❶ Bad news 1: when there is a few FLOP per iteration, the computation is bounded by memory transfer. Push/Pull to/from vector registers can kill your perfs!
- ❷ Bad news 2: the compiler can mispredict performance and it may not vectorize, even straightforward code.
- ❸ Good news: you can still sometimes be  $N$  times faster.  $N = \frac{\text{register size}}{\text{type size}}$

# Conclusion

- SIMD requires the proper data structure and fine-grained data parallelism
- Using vector registers and auto-vectorization, this can increase performance of data-parallel loop by a factor 4, 8, 16
- The compiler can auto-vectorize under *many* constraints, some code may be impossible to auto-vectorize
- The compiler can be wrong
- Debugging auto-vectorized code is **tedious**
- Be careful about the overhead for vectorization (push/pull to/from vector registers)

If you want something done right,  
you have to do it yourself!